


# Storage and computing resources vs Big Data

- How to store and process ever increasing amounts of data in a sustainable manner?
  - ♦ Both the data storage and computing resources should be scalable.
    - ▶ More data → more storage and computing resources.


- Approaches to scaling:
  - ♦ Scale up (vertical scaling).
    - ▶ Upgrade the computing nodes.
      - More RAM, better CPU, larger disk drives, etc.
  - ♦ Scale out (horizontal scaling):
    - ▶ Add more computing nodes.



(Image source: Centric Consulting)

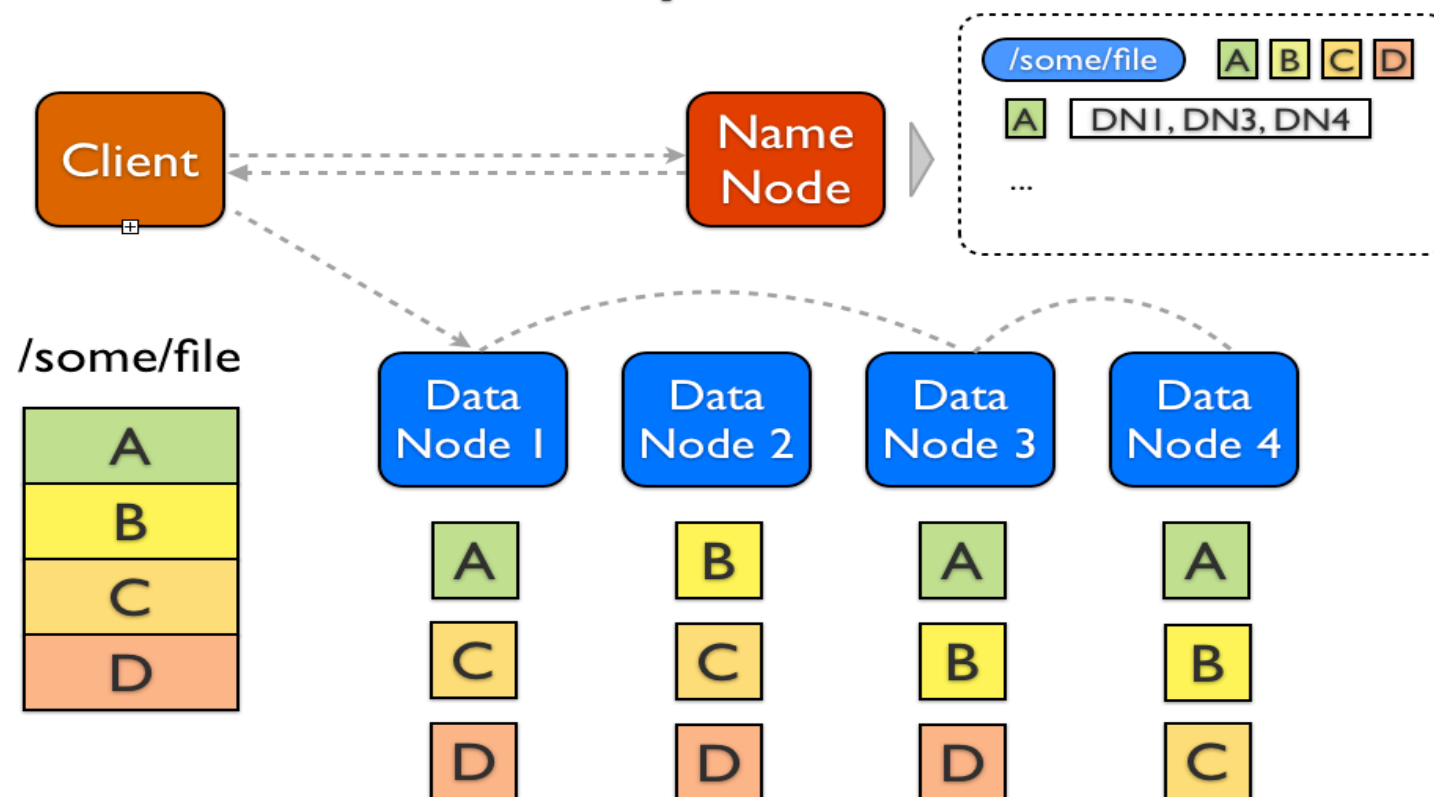
# Storage and computing resources vs Big Data...

- Vertical scaling:
  - ♦ Advantage: improvement is universal.
  - ♦ Disadvantage: limits of speed and scale of hardware improvement.
- Horizontal scaling → parallel / distributed storage and computation.
  - ♦ Advantages: in principle infinite scaling, can tolerate hardware errors.
  - ♦ Disadvantage: improvement is confined to parallelizable tasks.
- Big data: emphasis on horizontal scaling.
  - ♦ The ability to scale vertically is not enough for handling truly big data.



# Hadoop Distributed File System (HDFS)

- HDFS stores files in a cluster of name nodes and data nodes.
  - ♦ Data nodes: store the actual file data in distributed manner.
    - File is split into blocks, blocks stored on different data nodes.
    - Data blocks are also replicated in order to gain error tolerance.
  - ♦ Name nodes: file information (e.g. which data nodes store its blocks).



# Parallel computation

- Example 1: iterative computation of square root with  $n$ -decimal accuracy.
  - ♦ Note : the pseudocode is not truly in any real programming language...
  - ♦ The code does not parallelize well: strict dependence between steps.

```
// Assume that BigFloat can handle arbitrary-precision floats
BigFloat sqrt(BigFloat x, int n)
{
    low = 0
    high = max(1, x)
    // Assume that v.decimals(n) gives the
    // value of v rounded to n-decimal precision
    while(low.decimals(n) != high.decimals(n))
    { // Binary search until solution good enough
        mid = (low + high) / 2
        if(mid*mid < x)
            low = mid
        else
            high = mid
    }
    return low.decimals(n)
}
```


**Sqrt(3, 2) ≈ 1.73:**

Low: 0	High: 3
Low: 1.5	High: 3
Low: 1.5	High: 2.25
Low: 1.5	High: 1.875
Low: 1.688	High: 1.875
Low: 1.688	High: 1.782
Low: 1.688	High: 1.735
Low: 1.712	High: 1.735
Low: 1.724	High: 1.735
Low: 1.730	High: 1.735
Low: 1.730	High: 1.732

# Parallel computation...

- Example 2: computing moving average (the black line in the plot).
- ♦ Does parallelize well.

```
def m_avg(lst, i, m):  
    s = max(0, i-m/2)  
    e = min(len(lst), i + m/2)  
    total = 0  
    for i in range(s, e):  
        total += lst[i]  
    return total/(e-s)  
  
def moving_average(data, m):  
    result = []  
    for i in range(len(data)):  
        result.append(m_avg(data, i, m))  
    return result
```



# Parallel computation...

- Example 2: computing moving average.

- ♦ Does parallelize well: one simple possibility is sketched below.

```
def m_avg(lst, i, m):  
    s = max(0, i-m/2)  
    e = min(len(lst), i + m/2)  
    total = 0  
    for i in range(s, e):  
        total += lst[i]  
    return total/(e-s)
```

```
def slice_ma(data, m, a, b):  
    result = []  
    for i in range(a, b):  
        result.append(m_avg(data, i, m))  
    return result
```

```
def sliced_moving_average(data, m, slice_size):  
    result = []  
    i = 0  
    while i < len(data):  
        result += slice_ma(data, m, i, min(i+slice_size, len(data)))  
        i += slice_size  
    return result
```



These function calls could be run in parallel: they are independent of each other.

# Amdahl's law

- Amdahl's law:
  - ♦ Maximum parallel speedup using **n** computers:  $1 / (f + (1-f)/n)$ 
    - Here **f** is the fraction of code that can not be parallelized (is "serial").

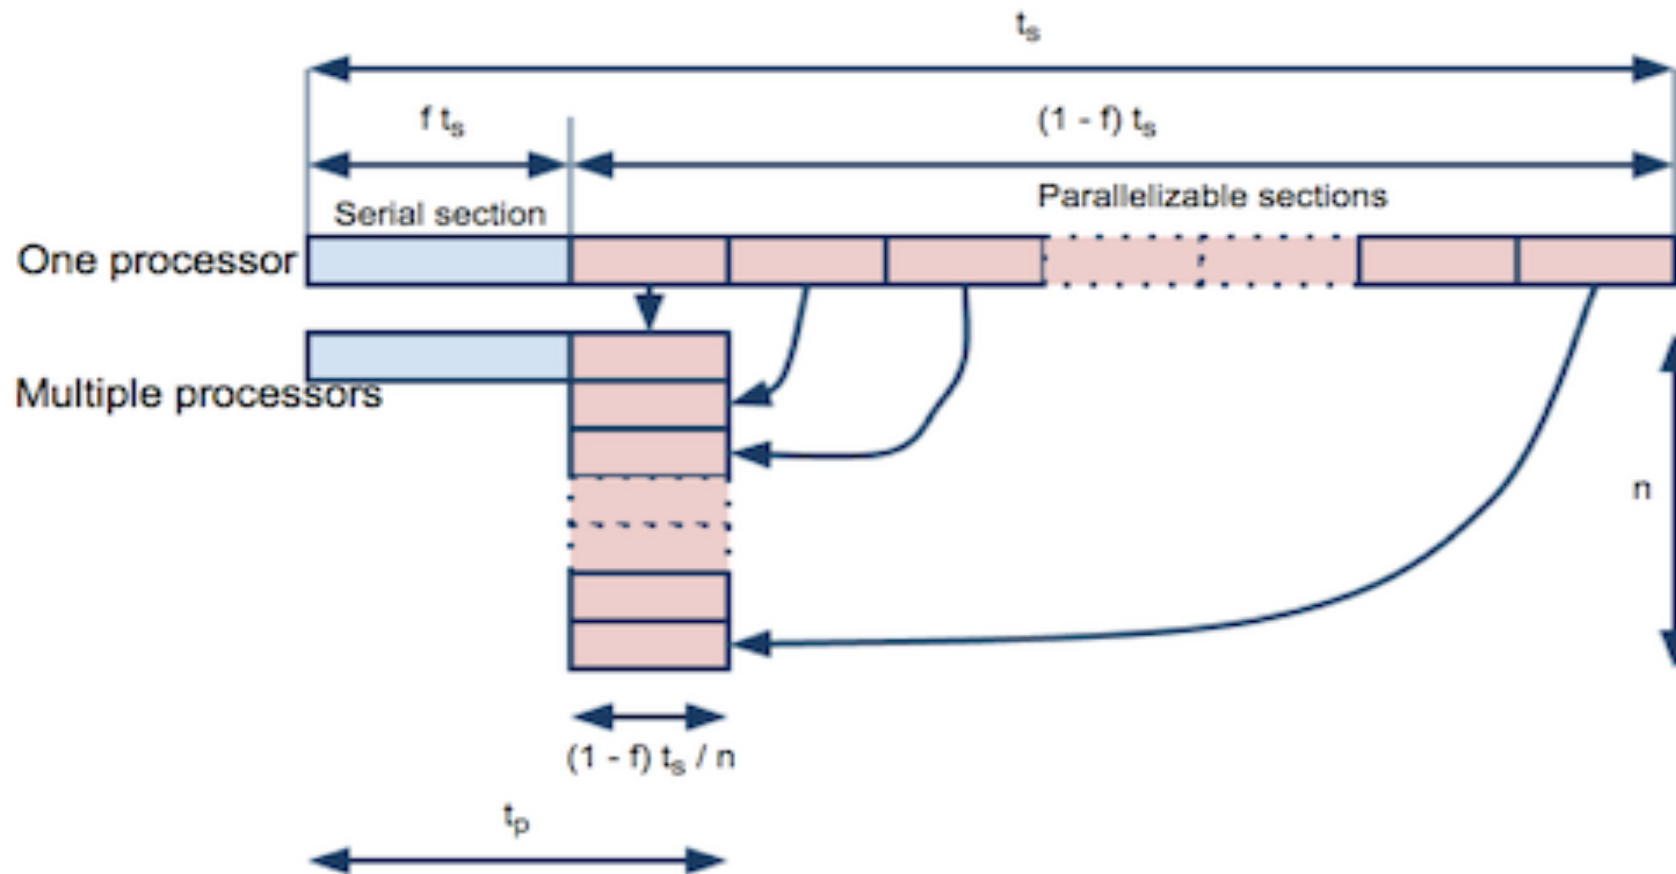


Figure source: Carlos P Sosa (IBM / Patton Fast Supercomputing Institute)

# Amdahl's law...

- Amdahl's law:
  - ♦ Maximum parallel speedup using **n** computers:  $1 / (\mathbf{f} + (1-\mathbf{f})/\mathbf{n})$ 
    - Here **f** is the fraction of code that can not be parallelized (is "serial").




Figure source: Carlos P Sosa (IBM / Patton Fast Supercomputing Institute)



# Hadoop MapReduce

- Parallel computation using two basic operations.
  - ♦ Map: group the data into chunks (that will be processed by same node).
  - ♦ Reduce: process chunks, output = the overall result (of one iteration).

## Hadoop MapReduce

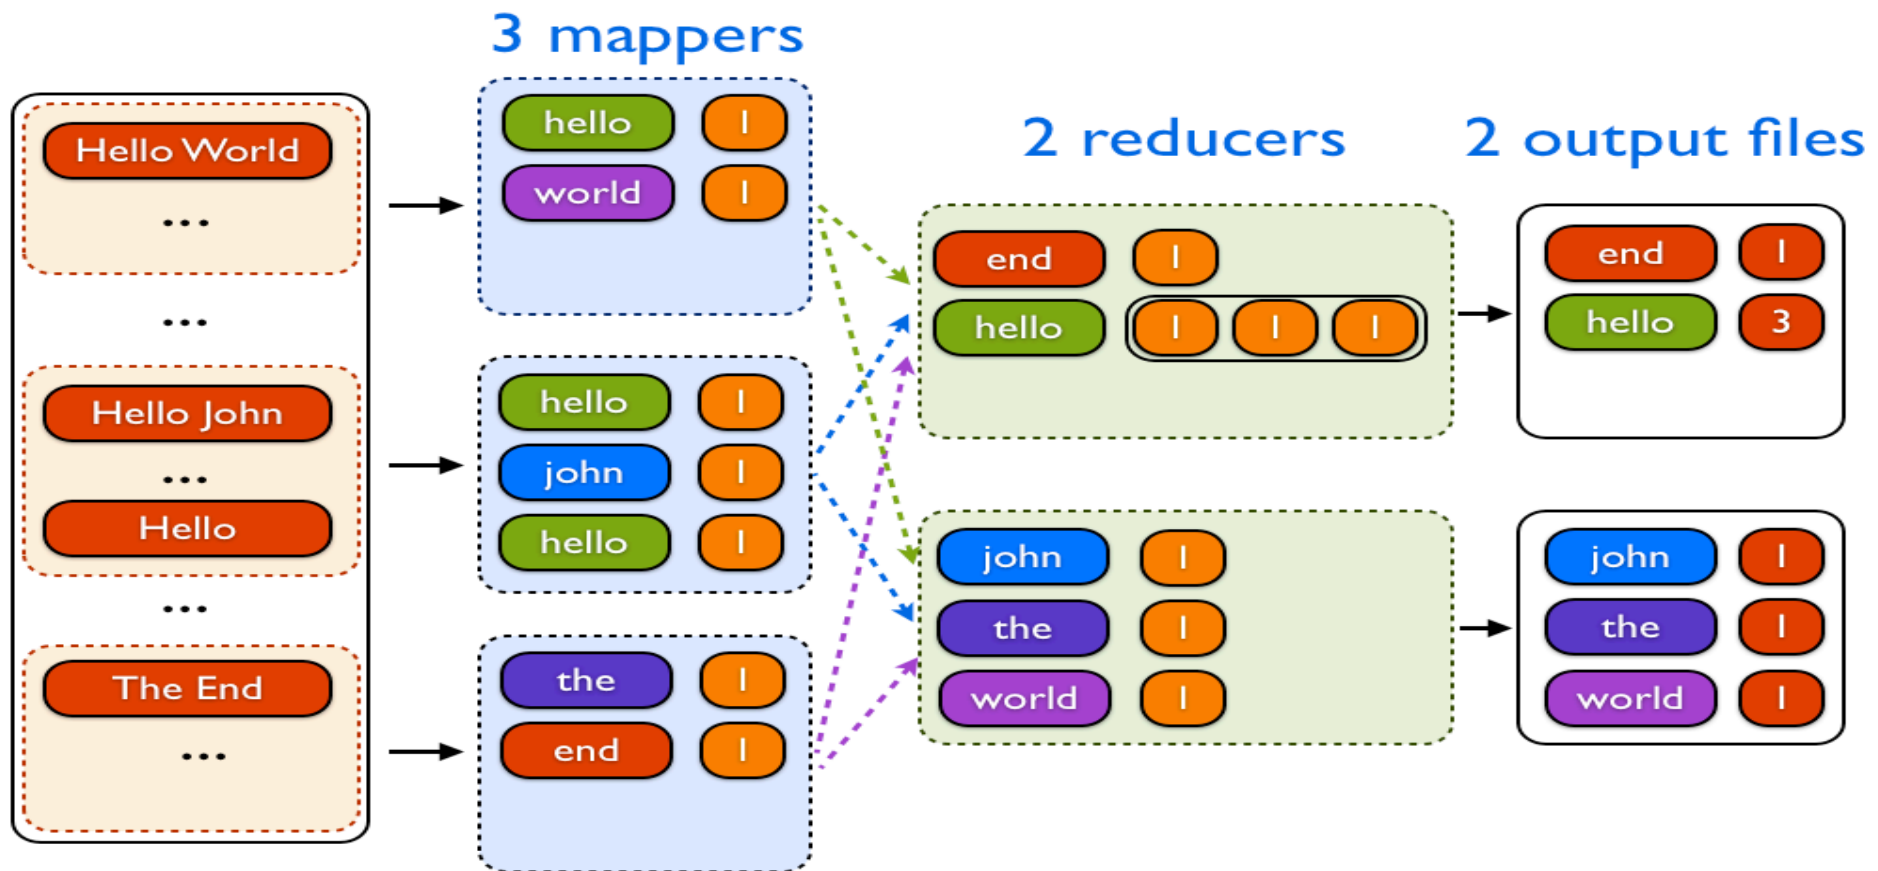


Figure source: Fernando Rodriguez Olivera / Nosqlessentials